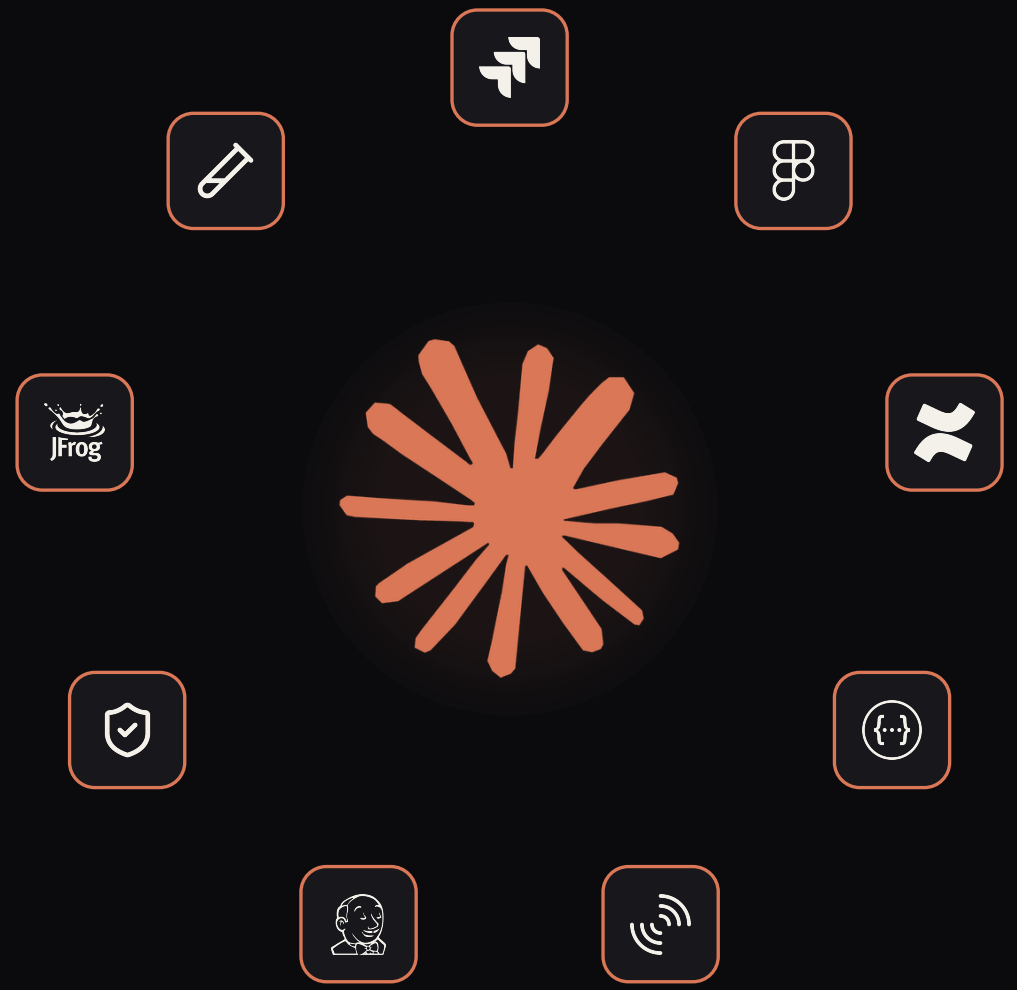
COFORGE · AMERIPRISE FINANCIAL

the-agentic-delivery-pipeline

Jira story in. Tested, scanned, review-ready code out.



Deepak Kumar · 22 October 2026

Built on Claude models, via GitHub Copilot and Kiro

WHO'S PRESENTING

Deepak Kumar

Coforge since 30 May 2024
At Ameriprise Financial from day one

Frontend team, 40–50 developers
The team using these agents today

Built the-agentic-delivery-pipeline
Solo, with some help from DevSecOps



Claude Certified Developer

Foundations · Anthropic · passed 9 Oct 2026

882

score
pass mark 720

QUESTION 01

What does a 13-point story cost us today?



What does a 13-point story cost us today?

40 man-hours, and most of it isn't writing code

40

man-hours

5 man-days for one 13-point story

Hunting for context

Jira, Figma, Confluence and SwaggerHub, one tab at a time

Chasing incomplete stories

Gaps surface mid-build, then wait on the reporter

Fixing scan findings

SonarQube, Veracode and Xray, one issue at a time

Waiting on a Jenkins expert

Read the logs, find someone, wait for their slot

Writing unit tests by hand

Slow, repetitive work to reach coverage

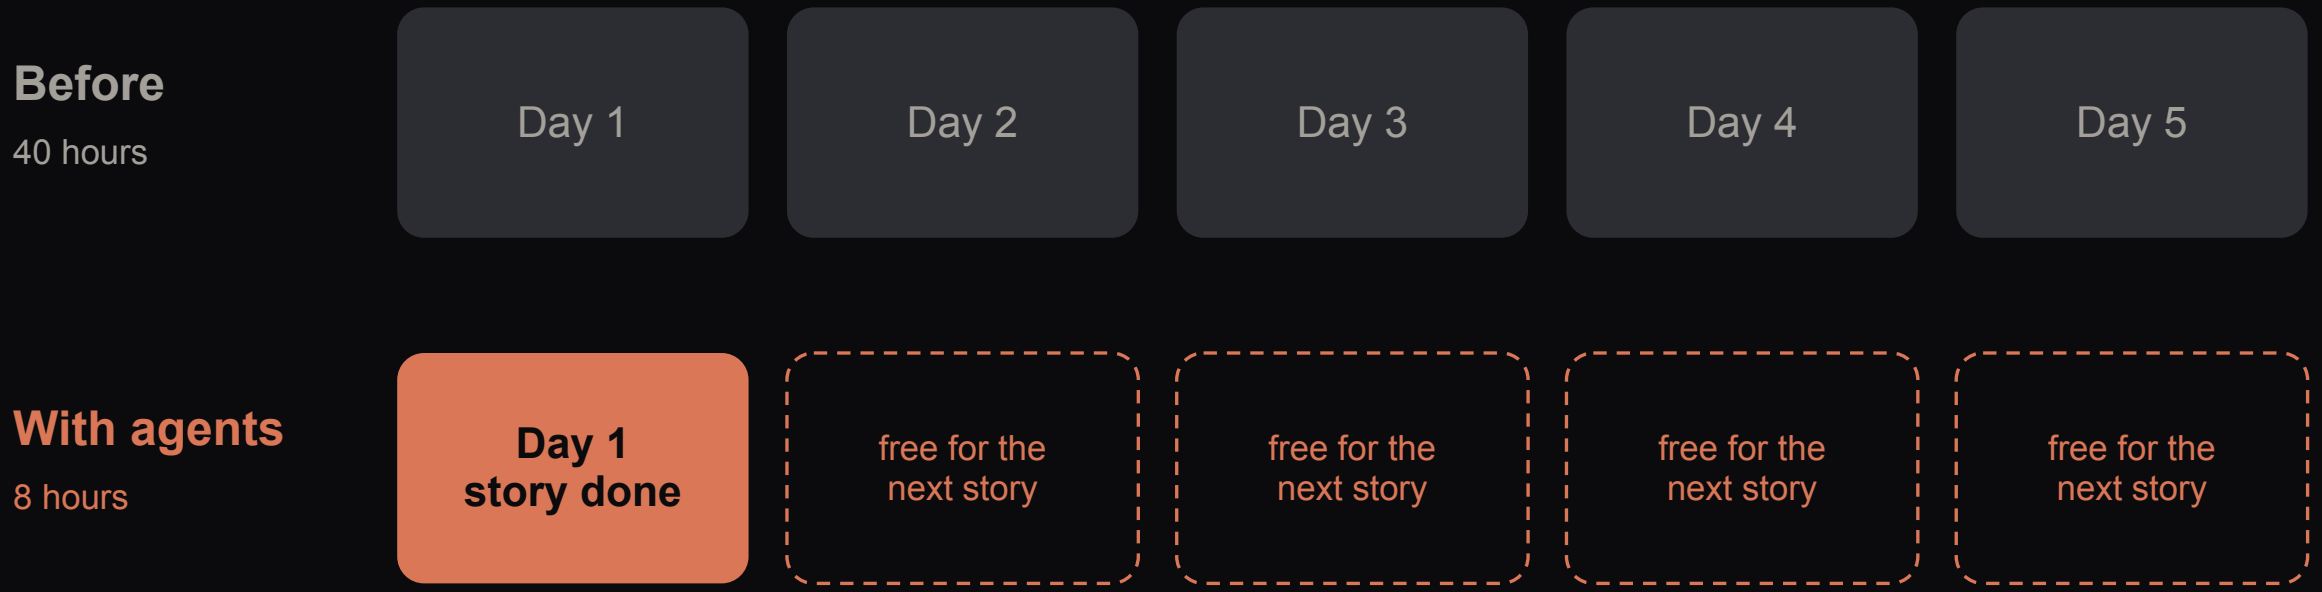
QUESTION 02

What if it took one day?



What if it took one day?

Same 13-point story. One man-day.



−80% *effort per story. The bigger the story, the bigger the saving.*

QUESTION 03

How?



How?

Nine atomic agents. Each does one job, end to end.

`jira-agent`

Jira

Reads the story, done and backlog stories, and the code

`figma-agent`

Figma MCP

Builds from the node ID with real library components

`swagger-hub-agent`

SwaggerHub

Writes API calls straight from the spec

`confluence-assistant`

Confluence

Pulls linked docs as precise context; publishes docs

`sonar-reporter`

SonarQube

Fetches issues, fixes them, rebuilds, repeats

`the-ci-analyzer`

Jenkins

Reads the console, fixes, commits, reruns

`veracode-analyser`

Veracode · SAST

Uploads, waits for the scan, fixes the flaws

`xray-analyzer`

JFrog Xray · SCA

Fixes dependency flaws through our own MCP

`code-coverage-agent`

Unit tests

Writes tests and loops until 95% coverage

Built solo · Claude models via GitHub Copilot and Kiro

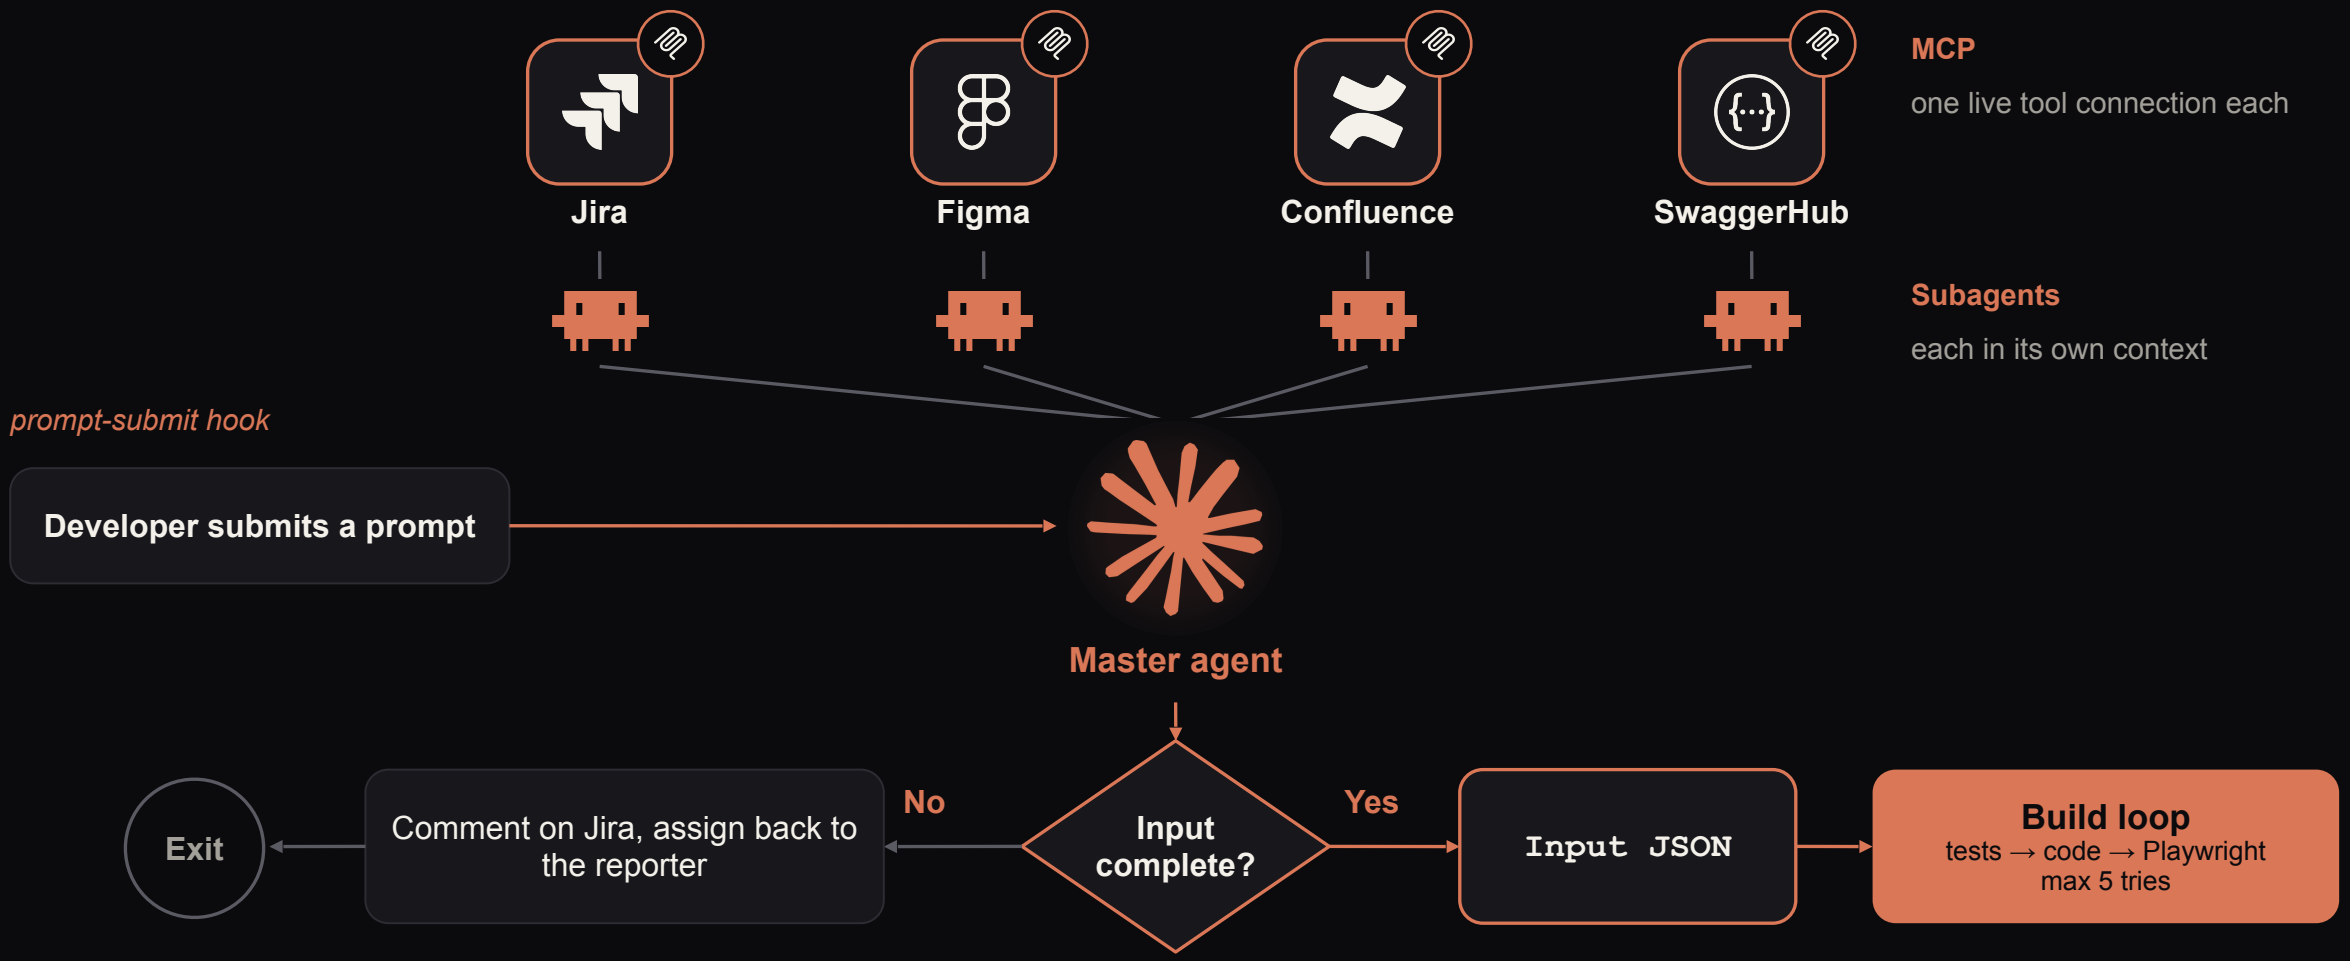
QUESTION 04

What happens when the agents work together?



What happens when the agents work together?

A master agent conducts. Specialists do the work.



QUESTION 05

What if the story itself is incomplete?



What if the story itself is incomplete?

It stops early and sends the story back



WEB-1234 · Account summary card



`jira-agent` commented just now

Before I build this, two things are missing:

- Acceptance criteria for the error state
- The Figma link for the mobile layout

Assigning back to the reporter.

Assignee → Reporter

Illustrative example



No tokens burned on guesswork

The run exits before any build work starts



No hours spent on the wrong build

The reporter fixes the story first



Checks more than one story

Done and backlog stories, plus the code itself

QUESTION 06

Who fixes the build when it breaks?



Who fixes the build when it breaks?

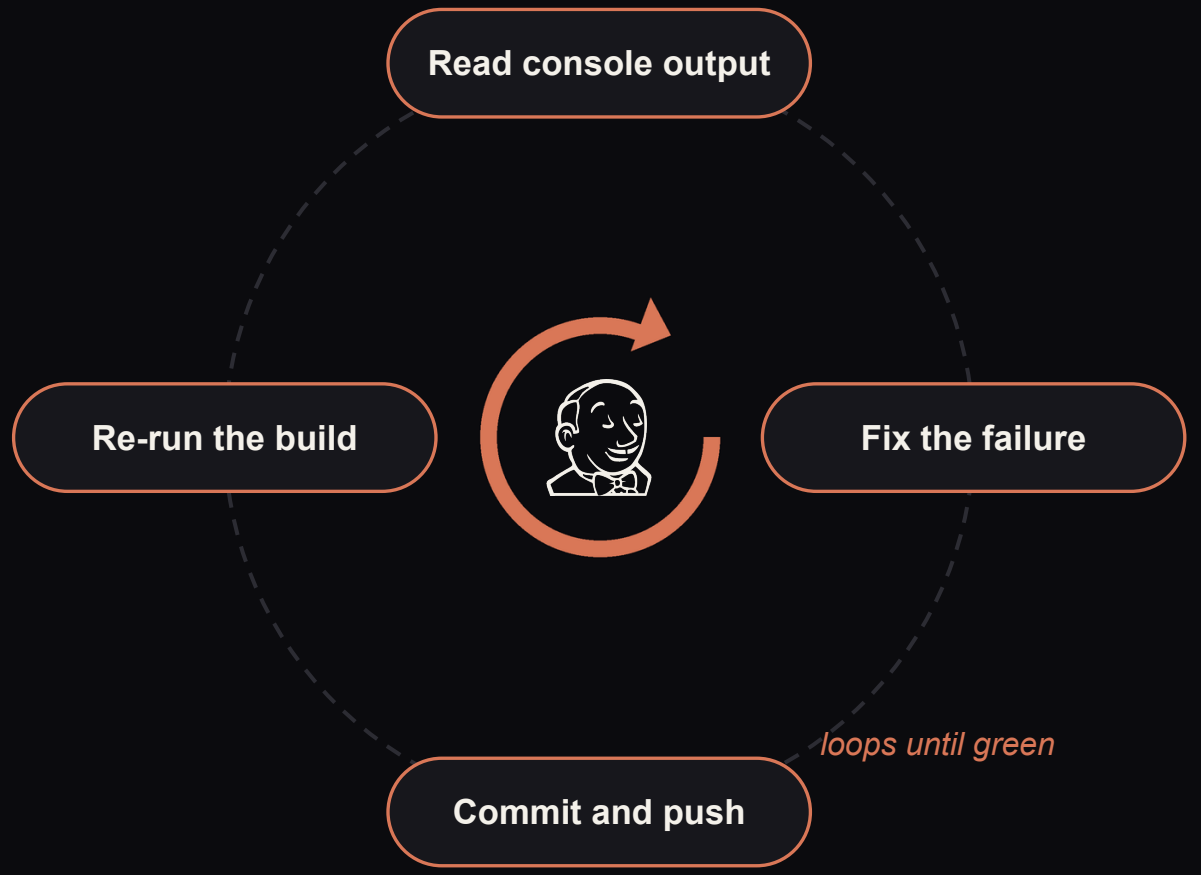
the-ci-analyzer fixes it and re-runs the build

BEFORE

- Log in to Jenkins
- Read the console logs
- Find a Jenkins expert
- Wait for their availability
- Apply the suggested fix

~90% less effort

NOW



QUESTION 07

And code quality and security?



And code quality and security?

Scan, fix, rebuild, rescan, until it's clean

SonarQube

Code quality

`sonar-reporter`



Fetch issues → fix
→ build → recheck

Veracode

SAST

`veracode-analyser`



Zip → upload → scan
→ fix → rescan

JFrog Xray

SCA

`xray-analyzer`



Read flaws via in-house MCP
→ fix → rescan

~90% less effort. The 10% left is PR review.

QUESTION 08

Two man-months of unit tests?



Two man-months of unit tests?

95% coverage in under two weeks, at a tenth of the cost

TIME

Human only

8 weeks · 2 man-months

`code-coverage-agent`

< 2 weeks



Writes tests, measures coverage, repeats until it reaches 95%

COST

~~$11,200~~

people: $35/hr × 8 hrs × 5 days × 8 weeks

~$1,120

AI credits on GitHub Copilot, or less

10× cheaper

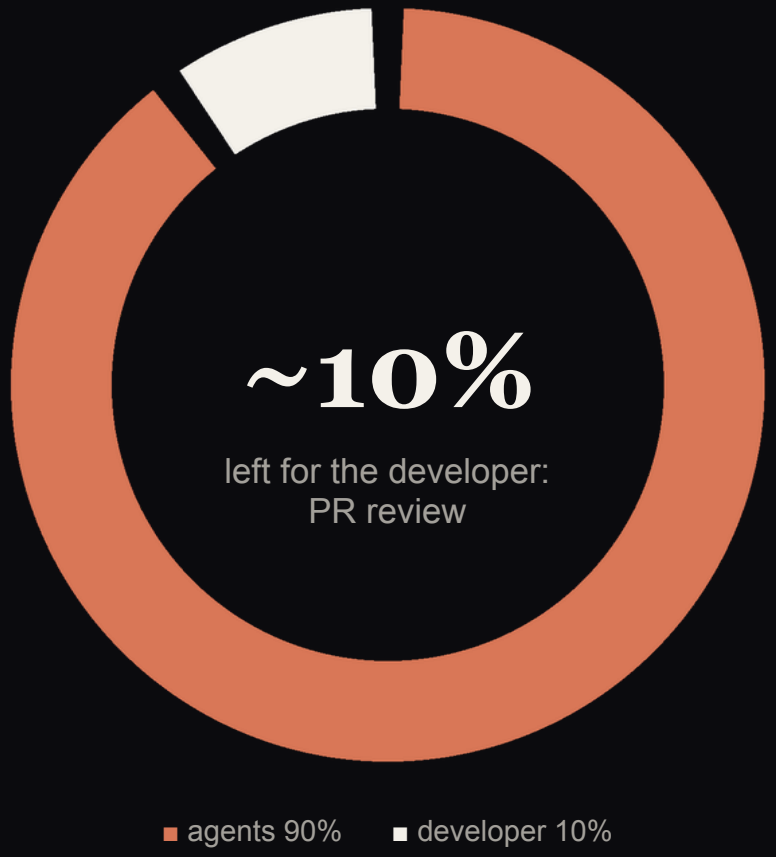
QUESTION 09

So what does the developer do now?



So what does the developer do now?

Developers move from writing code to reviewing it



BEFORE	NOW
~~Writes the code~~	Reviews the PR
~~Writes the tests~~ →	Makes the calls
~~Fixes scan findings~~	Owns quality
~~Debugs broken builds~~	

QUESTION 10

Doesn't AI get expensive?



Doesn't AI get expensive?

Tokens are spent only where they earn their keep

Fail fast
Incomplete stories exit before any build work

Prompt caching
Stable context is reused, not paid for again

Right-sized models
Cheaper models for subagents, set in each agent's definition

Trimmed tool output
Logs and scan results filtered before the model sees them

Never switch models mid-session
It throws away the cache and costs more

BUDGET TODAY

Monthly per-seat quota
GitHub Copilot AI credits, Kiro

$35 per dev seat / month
Figma MCP, dev seats only

Per-task token caps
Not available in Kiro yet

QUESTION 11

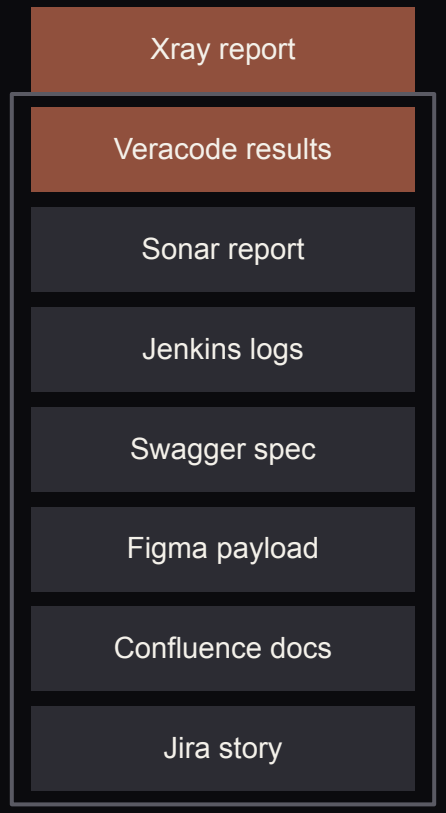
Won't the AI drown in context?



Won't the AI drown in context?

Every agent works in a small, clean context

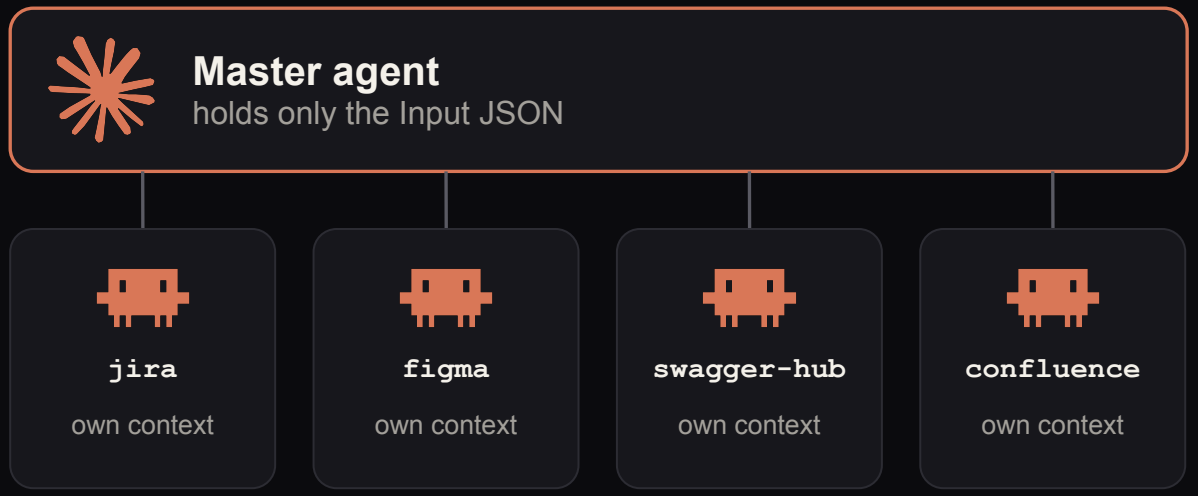
ONE AGENT DOING EVERYTHING



Overflow

Everything piles into one window. Quality drops, cost climbs.

MASTER + SPECIALISTS



- Subagents return results, not their whole context
- The Input JSON is the contract between steps
- Tool output is trimmed before it reaches the model

QUESTION 12

Is anyone actually using it?



Is anyone actually using it?

Live with Ameriprise's frontend team, and scaling



40–50 developers on the frontend team

Not yet all of Ameriprise. Scaling now.

40–50

developers using it

9

atomic agents, live

1

builder, designed and built solo

Live

Jira → UI pipeline, end to end

QUESTION 13

What's still slow?



What's still slow?

Next: other teams plug in through their own MCP servers

TODAY

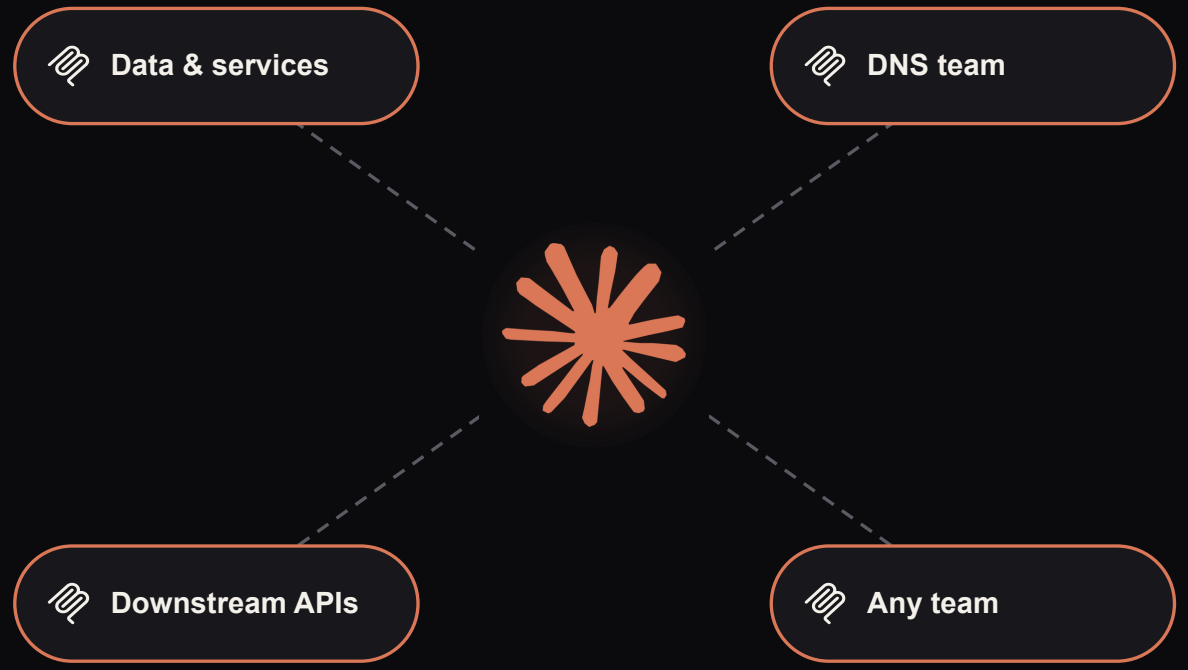
RE: RE: FW: timeout calling the accounts service

RE: FW: who owns this DNS record?

FW: still waiting on data & services

Long mail chains
while the SLA clock runs

NEXT



Each team publishes an MCP server; our agents work with it directly, with no mail chain.

QUESTION 14

What could Coforge do with this?



What could Coforge do with this?

Turn it into a Coforge accelerator for every account

1

Proven

Live at Ameriprise with a 40–50 developer frontend team

→

2

Packaged

Rebuilt as a Coforge-owned accelerator that fits any client

→

3

Offered

Sold to other accounts as a Coforge product

Why now

Hiring is nearly frozen. Clients expect more from the same teams.

I'd like to help build it.

Thank you

Questions?

Deepak Kumar · Coforge @ Ameriprise Financial

`the-agentic-delivery-pipeline`

